

Assignment-4

Singam Manasvi

Course Id:CS 5710

CRN :23921

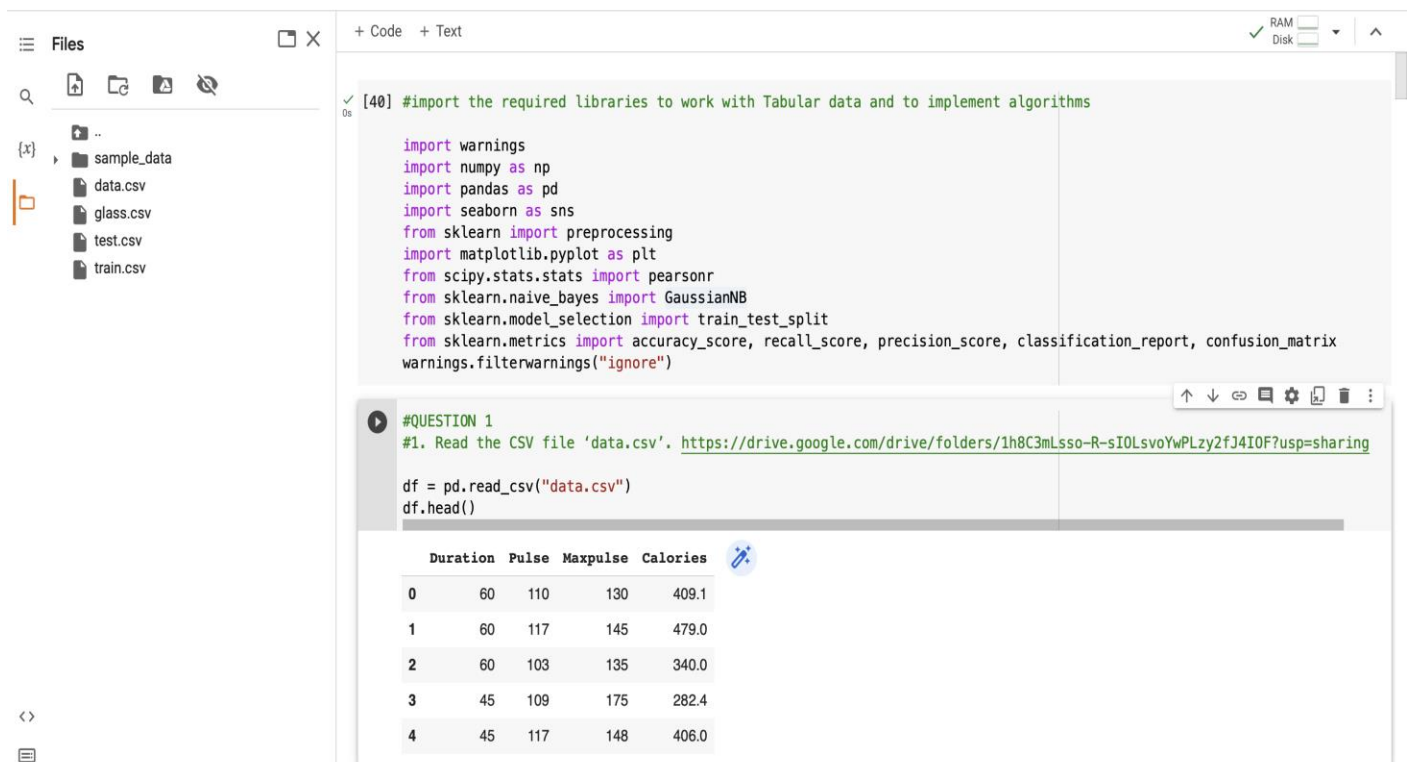
Student Id:700742501

Github Link:

<https://github.com/singammanasvi9440/Assignment-4>

1. Pandas

1.



The screenshot shows a Jupyter Notebook environment. On the left, a file explorer displays a directory named 'sample_data' containing four CSV files: 'data.csv', 'glass.csv', 'test.csv', and 'train.csv'. The main area contains two code cells. The first cell, labeled '[40]', imports various libraries including warnings, numpy, pandas, seaborn, sklearn preprocessing, matplotlib, scipy stats, and sklearn metrics. The second cell, labeled '#QUESTION 1', reads the 'data.csv' file from a Google Drive link and displays the first five rows of the data. Below the code, a preview table shows the data with columns: Duration, Pulse, Maxpulse, and Calories.

```
[40] #import the required libraries to work with Tabular data and to implement algorithms

import warnings
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn import preprocessing
import matplotlib.pyplot as plt
from scipy.stats.stats import pearsonr
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score, confusion_matrix
warnings.filterwarnings("ignore")
```

```
#QUESTION 1
#1. Read the CSV file 'data.csv'. https://drive.google.com/drive/folders/1h8C3mLsso-R-sI0LsvoYwPLzy2fJ4IOF?usp=sharing

df = pd.read_csv("data.csv")
df.head()
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0

Read the CSV file 'data.csv'.

2.

#2. Show the basic statistical description about the data.

```
df.describe()
```

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

Show the basic statistical description about the data using describe ().

3,4.

#3. Check if the data has null values.

```
df.isnull().any()
```

```
Duration    False
Pulse       False
Maxpulse    False
Calories    True
dtype: bool
```

#3a. Replace the null values with the mean

```
df.fillna(df.mean(), inplace=True)
df.isnull().any()
```

```
Duration    False
Pulse       False
Maxpulse    False
Calories    False
dtype: bool
```

#4. Select at least two columns and aggregate the data using: min, max, count, mean.

```
df.agg({'Maxpulse': ['min', 'max', 'count', 'mean'], 'Calories': ['min', 'max', 'count', 'mean']})
```

	Maxpulse	Calories
min	100.000000	50.300000
max	184.000000	1860.400000
count	169.000000	169.000000
mean	134.047337	375.790244

Check if the data has null values using isnull().

Replace the null values with the df.mean().

Select at least two columns and aggregate the data using: min, max, count, mean.

5.

```
#5. Filter the dataframe to select the rows with calories values between 500 and 1000.
```

```
df.loc[(df['Calories']>500)&(df['Calories']<1000)]
```

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

Filter the data frame as per the requirements.

6,7.

```
#6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
```

```
df.loc[(df['Calories']>500)&(df['Pulse']<100)]
```

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

```
#7. Create a new "df_modified" dataframe that contains all the columns from df except for "Maxpulse".
```

```
df_modified = df[['Duration','Pulse','Calories']]  
df_modified.head()
```

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

Filter the data frame and create a new modified data frame.

8,9.

```
[11] #8. Delete the "Maxpulse" column from the main df dataframe
```

```
del df['Maxpulse']
```

```
df.head()
```

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

```
[13] df.dtypes
```

```
Duration      int64
Pulse         int64
Calories      float64
dtype: object
```

```
#9. Convert the datatype of Calories column to int datatype.
```

```
df['Calories'] = df['Calories'].astype(np.int64)
df.dtypes
```

```
Duration      int64
Pulse         int64
Calories      int64
dtype: object
```

Delete the maxpulse using del


Convert the datatype of calories column into int datatype.

10.

```
#10. Using pandas create a scatter plot for the two columns (Duration and Calories).
```

```
df.plot.scatter(x='Duration',y='Calories',c='blue')
```

```
<Axes: xlabel='Duration', ylabel='Calories'>
```



Scatter plot for duration and Calories.

1. Titanic Dataset

1.

```
#1. TITANIC DATASET
Loading the data file into the program
df=pd.read_csv("train.csv")

df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
#converted categorical data to numerical values for correlation calculation

label_encoder = preprocessing.LabelEncoder()
df['Sex'] = label_encoder.fit_transform(df.Sex.values)

#Calculation of correlation for 'Survived' and 'Sex' in data
correlation_Value= df['Survived'].corr(df['Sex'])

print(correlation_Value)
```

[-0.5433513806577555]

Ans: Yes, we should keep the 'Survived' and 'Sex' features helps classify the data accurately

Finding correlation between 'survived' (target column) and 'sex' column.

```
#print correlation matrix
matrix = df.corr()
print(matrix)
```

```

PassengerId  PassengerId  Survived  Pclass  Sex  Age  SibSp  \
PassengerId  1.000000 -0.005007 -0.035144  0.042939  0.036847 -0.057527
Survived     -0.005007  1.000000 -0.338481 -0.543351 -0.077221 -0.035322
Pclass       -0.035144 -0.338481  1.000000  0.131900 -0.369226  0.083081
Sex          0.042939 -0.543351  0.131900  1.000000  0.093254 -0.114631
Age          0.036847 -0.077221 -0.369226  0.093254  1.000000 -0.308247
SibSp        -0.057527 -0.035322  0.083081 -0.114631 -0.308247  1.000000
Parch        -0.001652  0.081629  0.018443 -0.245489 -0.189119  0.414838
Fare         0.012658  0.257307 -0.549500 -0.182333  0.096067  0.159651

PassengerId  Parch  Fare
PassengerId -0.001652  0.012658
Survived     0.081629  0.257307
Pclass       0.018443 -0.549500
Sex          -0.245489 -0.182333
Age          -0.189119  0.096067
SibSp        0.414838  0.159651
Parch        1.000000  0.216225
Fare         0.216225  1.000000
```

2.


```
# One way of visualizing correlation matrix in form of spread chart
```

```
df.corr().style.background_gradient(cmap="Reds")
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159651
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	0.216225	1.000000

#Second form of visualizing correlation matrix using heatmap() from seaborn

```
sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



Two visualizations to show correlations.

3.

#Loaded data files test and train and merged files

```
train_raw = pd.read_csv('train.csv')
test_raw = pd.read_csv('test.csv')
train_raw['train'] = 1
test_raw['train'] = 0
df = train_raw.append(test_raw, sort=False)
features = ['Age', 'Embarked', 'Fare', 'Parch', 'Pclass', 'Sex', 'SibSp']
target = 'Survived'
df = df[features + [target] + ['train']]
df['Sex'] = df['Sex'].replace(['female', 'male'], [0, 1])
df['Embarked'] = df['Embarked'].replace(['S', 'C', 'Q'], [1, 2, 3])
train = df.query('train == 1')
test = df.query('train == 0')
```

Drop missing values from the train set.

```
train.dropna(axis=0, inplace=True)
labels = train[target].values
train.drop(['train', target, 'Pclass'], axis=1, inplace=True)
test.drop(['train', target, 'Pclass'], axis=1, inplace=True)
```

#Test and train split

```
X_train, X_val, Y_train, Y_val = train_test_split(train, labels, test_size=0.2, random_state=1)
```

```
classifier = GaussianNB()
```

```
classifier.fit(X_train, Y_train)
```

▼ GaussianNB
GaussianNB()


```

y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(Y_val, y_pred))

```

	precision	recall	f1-score	support
0.0	0.79	0.80	0.80	85
1.0	0.70	0.69	0.70	58
accuracy			0.76	143
macro avg	0.75	0.74	0.75	143
weighted avg	0.75	0.76	0.75	143

```

[[68 17]
 [18 40]]
accuracy is 0.7552447552447552

```

Implementing Naïve Bayes method using scikit-learn library and calculate the accuracy.

2.GLASS DATASET

```

glass=pd.read_csv("../glass.csv")
glass.head()

```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0	1
2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1


```

glass.corr().style.background_gradient(cmap="Reds")

```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
RI	1.000000	-0.191885	-0.122274	-0.407326	-0.542052	-0.289833	0.810403	-0.000386	0.143010	-0.164237
Na	-0.191885	1.000000	-0.273732	0.156794	-0.069809	-0.266087	-0.275442	0.326603	-0.241346	0.502898
Mg	-0.122274	-0.273732	1.000000	-0.481799	-0.165927	0.005396	-0.443750	-0.492262	0.083060	-0.744993
Al	-0.407326	0.156794	-0.481799	1.000000	-0.005524	0.325958	-0.259592	0.479404	-0.074402	0.598829
Si	-0.542052	-0.069809	-0.165927	-0.005524	1.000000	-0.193331	-0.208732	-0.102151	-0.094201	0.151565
K	-0.289833	-0.266087	0.005396	0.325958	-0.193331	1.000000	-0.317836	-0.042618	-0.007719	-0.010054
Ca	0.810403	-0.275442	-0.443750	-0.259592	-0.208732	-0.317836	1.000000	-0.112841	0.124968	0.000952
Ba	-0.000386	0.326603	-0.492262	0.479404	-0.102151	-0.042618	-0.112841	1.000000	-0.058692	0.575161
Fe	0.143010	-0.241346	0.083060	-0.074402	-0.094201	-0.007719	0.124968	-0.058692	1.000000	-0.188278
Type	-0.164237	0.502898	-0.744993	0.598829	0.151565	-0.010054	0.000952	0.575161	-0.188278	1.000000


```
sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



```
X_train, X_val, Y_train, Y_val = train_test_split(glass[:-1], glass['Type'], test_size=0.2, random_state=1)
classifier = GaussianNB()
classifier.fit(X_train, Y_train)
y_pred = classifier.predict(X_val)
# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
print('accuracy is', accuracy_score(Y_val, y_pred))
```

```

              precision    recall  f1-score   support

     1         0.90        0.95        0.92         19
     2         0.92        0.92        0.92         12
     3         1.00        0.50        0.67          6
     5         0.00        0.00        0.00          1
     6         1.00        1.00        1.00          1
     7         0.75        0.75        0.75          4

 accuracy          0.84
macro avg          0.76        0.69        0.71         43
weighted avg          0.89        0.84        0.85         43

[[18  1  0  0  0  0]
 [ 1 11  0  0  0  0]
 [ 1  0  3  2  0  0]
 [ 0  0  0  0  0  1]
 [ 0  0  0  0  1  0]
 [ 0  0  0  1  0  3]]
accuracy is 0.8372093023255814
```

```

from sklearn.svm import SVC, LinearSVC

classifier = LinearSVC()

classifier.fit(X_train, Y_train)

y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(Y_val, y_pred))

```

	precision	recall	f1-score	support
1	1.00	0.95	0.97	19
2	0.57	1.00	0.73	12
3	0.00	0.00	0.00	6
5	0.00	0.00	0.00	1
6	0.00	0.00	0.00	1
7	0.75	0.75	0.75	4
accuracy			0.77	43
macro avg	0.39	0.45	0.41	43
weighted avg	0.67	0.77	0.70	43

```

[[18  1  0  0  0  0]
 [ 0 12  0  0  0  0]
 [ 0  5  0  0  0  1]
 [ 0  1  0  0  0  0]
 [ 0  1  0  0  0  0]
 [ 0  1  0  0  0  3]]
accuracy is 0.7674418604651163

```

We got better accuracy for Naïve Bayes method which is 0.8372093023255814. Naive Bayes analysis works well with probabilistic concepts whereas Linear SVM works better with linear regression logics. But to perform more accurately SVM requires large amounts of data to train and test the data. So, due to the amount of data Naive Bayes algorithm gives better accuracy compared to Linear SVM.